

# **AWS Cloud Practitioner Training**

## **Module 05**

### **Serverless Solutions with Lambda**

Instructor: Tim Platt, Cloud Solution Architect

# DynamoDB

## DynamoDB

A Non-relational (NoSQL) database – Internet scale and accessed as a WEB SERVICE.



## Key Points

- Based on Key/Value access pattern (Does that sound familiar?)
- No upper limit on table size or number of rows (called “Items” in DynamoDB)
- Massive, internet scale read/write throughput (if you need it)
- Significantly different than relational database. You don’t use SQL. DynamoDB has its own Application Programming Interface (API) with GETs, PUTs, QUERIES, SCANS. Accessed over the Internet as a WEB SERVICE.
- Originally invented for the Amazon.com Shopping Cart
- SUPERPOWER: Schema-less database – great for varied, diverse datasets and JSON
- SUPERPOWER: Single digit millisecond latency for GETs and PUTs (single Item access using a key) at any scale
- SUPERPOWER: Managed service using multiple AZs automatically. No upgrades or maintenance to deal with.

**NOTE: Due to the extreme differences – migrating from Relational database to DynamoDB is a lot of work**

Let's create a DynamoDB  
Table to see how it is  
different than SQL  
databases...

# AWS Lambda

## Serverless Compute

Run code on demand without  
managing servers



## Key Points

- Upload function code (C#, Java, Python, JavaScript, etc.) and it can be triggered (run) on demand by the Lambda service.
- Scales automatically with no ASG or Load Balancer needed
- SUPERPOWER: Only pay while the code is running! Very efficient for sporadic or highly variable loads
- SUPERPOWER: Make it a REST API by using API Gateway or trigger it when an object is created in an S3 bucket, or when a message hits a SQS queue, or an SNS topic – there's many, many ways to trigger the Lambda to run.

BUT ...


- Function code can only run for up to 15 minutes
- The execution environment is ephemeral (temporary) – your code must be stateless (store state in a database or S3)
- COLD STARTs happen when the Lambda service doesn't have a previously used "WARM" environment and there's a little bit of delay (not much) before your code starts running

Let's create a simple  
Lambda function to  
understand how it works

# Simple Notification Service (SNS)

## Notifications

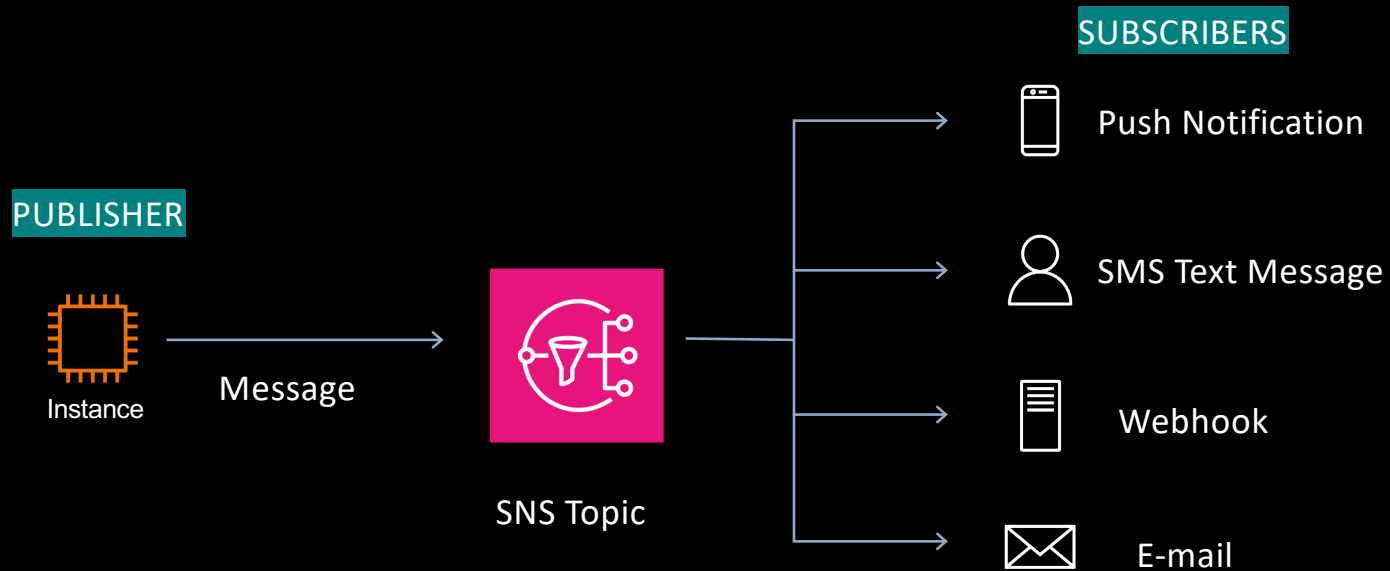
Send notifications using Publish and Subscribe (Pub-Sub) paradigm



## Key Points

- Your application or workflow can PUBLISH a message to a TOPIC and any and all SUBSCRIBERS will receive a notification
- Sometimes also known as a “fan-out” because ONE publisher can notify MANY subscribers.
- The messages are “pushed” to the subscribers – they don’t need to poll constantly.
- Publisher is “decoupled” from the Subscribers and needs no knowledge of the subscribers
- Very simple API to use
- SUPERPOWER: Very scalable – great for sending many, many, many messages to many, many, many subscribers
- SUPERPOWER: Send emails, text/SMS messages, push notifications – or invoke Lambda and other things
- SUPERPOWER: Use SNS topics to build simple event driven architectures. For example, an SNS Topic might push a notification to a Lambda function


# SNS Architecture – Fan Out (One to Many)



# Simple Queue Service (SQS)

## Queueing

Use a highly scalable QUEUE to buffer or store messages until you can process them




## Key Points

- A durable queue that you can use to temporarily store messages (like incoming customer orders) until you can process them.
- One or more PRODUCERS will put MESSAGES into the QUEUE. One or more CONSUMERS will process the messages – and delete them from the queue when they are done.
- The consumers will POLL to find any unprocessed messages in the queue.
- There are two basic types of Queue : Standard (AT LEAST ONCE DELIVERY) and FIFO – First In First Out – which guarantees delivery ONCE and in ORDER that the messages arrived.
- Standard queue scales much higher – but you must deal with potential message duplication
- SUPERPOWER: Very scalable and highly reliable
- SUPERPOWER: SQS queues can hold messages for up to 14 days.



# SQS Architecture



PRODUCER(S) are constantly polling, checking for messages in the queue. They process each one, then delete it from the queue.

# SNS Versus SQS

They are two very different things – even though both deal with "Messages"

## Persistence

SNS is not a buffer – there's no storage – the messages are sent immediately and not retained.

SQS can store messages for up to 14 days – allowing time for our consumers to get caught up

## Push vs Poll

SNS – Messages are PUSHed to the subscribers when they arrive

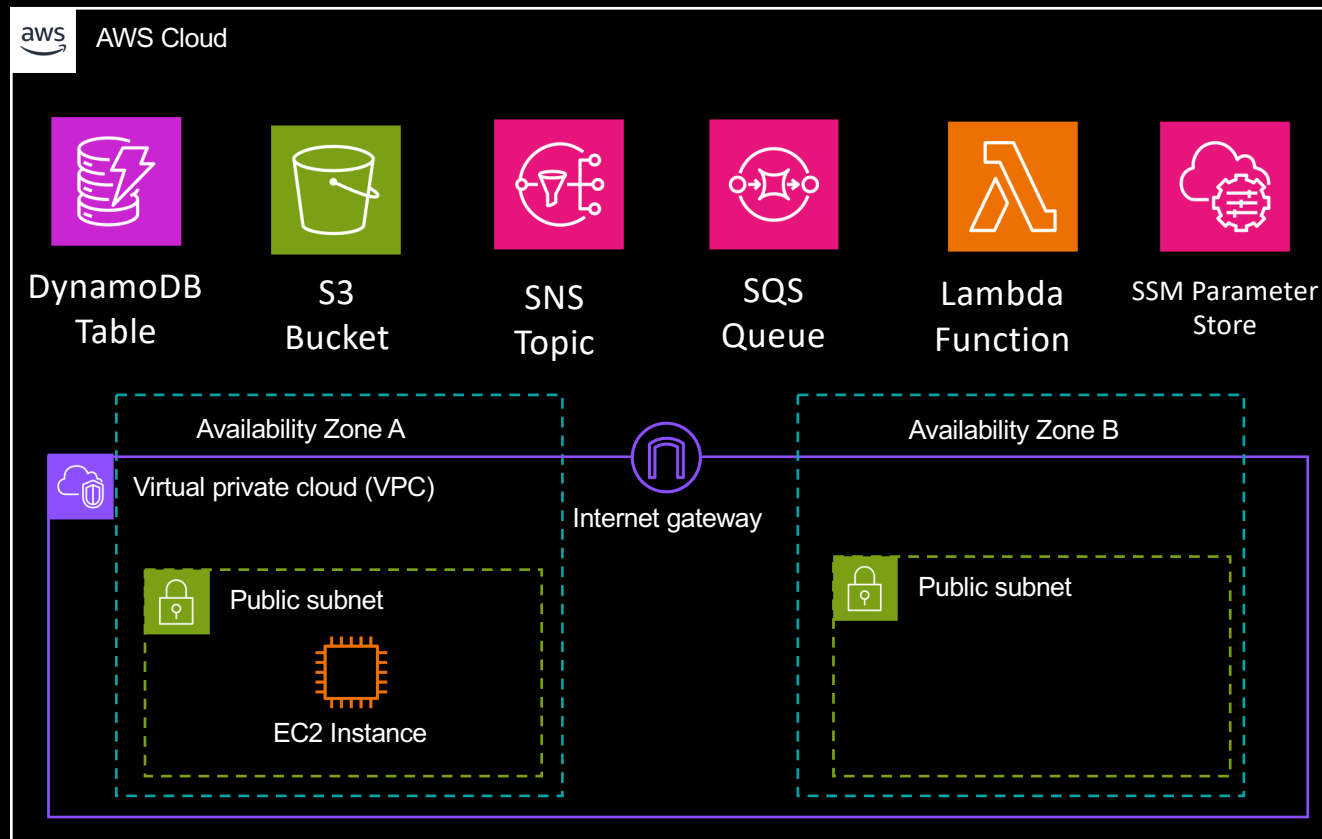
SQS – The consumers must poll repeatedly (in a loop) to constantly look for new messages

## Distribution


SNS Messages can go to many, many subscribers (One to Many)

SQS Messages are intended to be "consumed" by a **SINGLE** Consumer process (Imagine processing a customer order, for example)

# Serverless Services – live in the REGION, not your VPC



# Application Architecture



# Links

- Simple Notification Service (SNS) Developer Guide: <https://docs.aws.amazon.com/sns/latest/dg/welcome.html>
- Simple Queue Service (SQS) Developer Guide: <https://docs.aws.amazon.com/AWSSimpleQueueService/latest/SQSDeveloperGuide/welcome.html>
- AWS Lambda Developer Guide: <https://docs.aws.amazon.com/lambda/latest/dg/welcome.html>
- Lambda Execution environment lifecycle: <https://docs.aws.amazon.com/lambda/latest/dg/lambda-runtime-environment.html>
- Lambda Example – File Processing: <https://docs.aws.amazon.com/lambda/latest/dg/file-processing-app.html>
- Lambda Example – Cron Job: <https://docs.aws.amazon.com/lambda/latest/dg/scheduled-task-app.html>